

Coding: Buffer Overflows

David Basin

Department of Computer Science
ETH Zurich

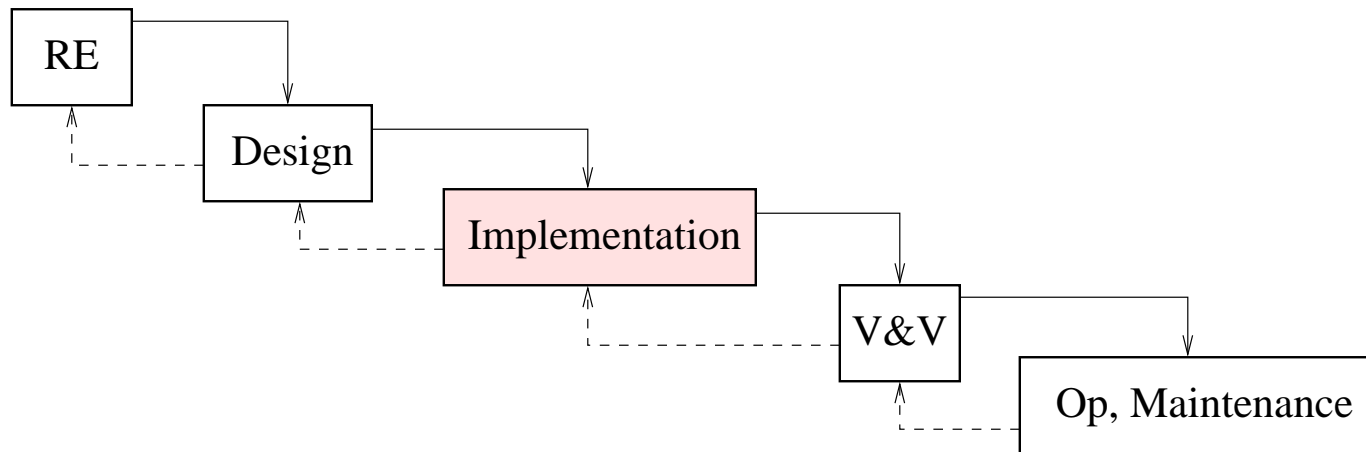
Road map¹

Motivation

- Pointers in C
- Compilation and memory layout
- Buffer overflows
- Defense
- Summary

¹Some of the material from this talk was developed by Heiko Mantel, for Security Engineering 2004/05

Context: implementation



It's possible to have a great design but an insecure implementation

1. Implementation deviates from design

2. Design leaves many options open (including unsafe ones)

E.g., design doesn't specify password lengths, their storage, ...

3. Concretization may introduce new vulnerabilities

E.g., buffer overflows and other attacks on data and control

Motivation — buffer overflows

- Public enemy #1

See e.g., CERT advisories or SANS top 20 vulnerabilities ranking

- Example: Morris Internet Worm (November 1988)

- ▶ Attacked Sun3 and VAX systems running BSD
- ▶ Fairly clever, e.g., named itself “sh” (hard to detect) and set maximum core-dump size to 0 bytes (hard to catch)
- ▶ Propagated itself using various exploits, including a buffer overflow attack against the finger daemon *fingerd*.
- ▶ Did not delete files, but resulting financial loss between \$100,000 and \$1,000,000 (US General Accounting Office)

Role of studying vulnerabilities (and a warning)

- Understanding vulnerabilities is vital to developing and evaluating countermeasures. Failure to consider them leads to insecure code or overconfidence in partial solutions
- Exploiting vulnerabilities (= hacking) is a criminal offense
- We expect you to use this knowledge in a responsible and ethical way
- Any experimentation must be done in a controlled environment

Buffer overflows

- A **buffer** is a contiguous region of memory storing data of the same type, e.g., characters
- A **buffer overflow** occurs when data is written past a buffer's end
- The resulting damage depends on
 - ▶ Where the data spills over to
 - ▶ How this memory region is used (e.g., flags for access control)
 - ▶ What modifications are made
- Example: Morris attack on finger
 - ▶ Intended use: `finger basin@inf.ethz.ch`
 - ▶ Abuse: `finger <exploit-code> ... <return-address>`
 - ▶ Overwrites the return address and provides exploit code

Road map

- Motivation

Pointers in C

- Compilation and memory layout
- Buffer overflows
- Defense
- Summary

C primer

- C (and C++) is powerful, widespread, and undisciplined!
- Basic data types
 - Char (1 byte), int (≥ 2 bytes), long (≥ 4 bytes) ...
- Pointers: expressions like `int *ptr`
 - ▶ `ptr` is the **address** of a memory cell
 - ▶ `*ptr` is the **contents** of the memory cell
- Address taken using `&`. For `i` and integer and `f` a function,
 - ▶ `&i` is the **address** of the cell storing `i`
 - ▶ `&f` is the **address** of a function `f`

Pointers — examples

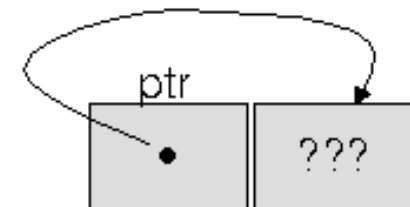

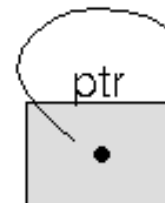

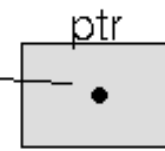

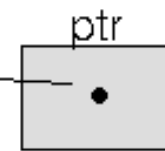

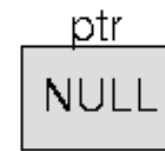

```
int i = 0;  
int *ptr = NULL;
```

```
ptr = &i;
```

```
*ptr = 15;
```

```
ptr = &ptr;
```

```
ptr++;
```



Arrays

- Declaring and accessing an array
 - ▶ `char buf[8];` declares a buffer for up to 8 characters
 - ▶ Elements are `buf[0]`, `buf[1]`, .., `buf[7]`
 - ▶ No bounds checking, so `buf[13]` is also possible
- Arrays are similar to pointers: store address where buffer begins
- Some equalities

<code>&(buf[0]) = buf</code>	<code>buf[0] = *(buf+0)</code>
<code>&(buf[1]) = buf + 1</code>	<code>buf[1] = *(buf+1)</code>

- Strings are arrays of characters, with `'\0'` at end
 - ▶ E.g., string "hi" represented by `'h'` `'i'` `'\0'`
 - ▶ No explicit information about length of strings
 - ▶ This savings is part of the problem!

Input, output and string handling functions

- `getchar()` reads a character from standard input
- `putchar(c)` writes a character to standard output
- `printf("Page %i has title %s\n", pageno, title)`
 - ▶ String printed to standard output may contain place holders replaced by arguments
 - ▶ Place holders provide formatting instructions.
E.g., `%i` stands for an integer in decimal and `%s` for a string
- Various other functions store strings in buffers. E.g.
 - ▶ `gets(dst)` read string from `stdin` into `dst`
 - ▶ `strcpy(dst,src)` copy string `src` into buffer `dst`
 - ▶ `sprintf(dst, <FmtStr>,<Exp>)` print string into buffer
 - ▶ `scanf, fscanf, ...`

A vulnerable program

buffer of
limited size

is written
without limits

```
#include <stdio.h>
int main()
{
    char buf[8]           // buffer for storing the input string
    char ch;              // auxiliary variable
    char *ptr = buf;       // auxiliary pointer
    while ((ch=getchar()) != '\n' // terminate on EOL
           && ch != -1)      // terminate on error
    {
        *ptr = ch;          // store character
        ptr++;              // increment pointer
    }
    *ptr = '\0';           // terminate the string
    printf("%s\n",buf);
    return 0;
}
```

Some example runs

```
$ ./my-getstring
example
example
```

```
$ ./my-getstring
long example
long example
Illegal instruction
```

per esempio >7 chars
ptr va a puntare (forse)
char ch, sovrascrivendolo

```
$ ./my-getstring
very long example
very long example
Segmentation fault
```

program tried to access outside the
memory image of its own process

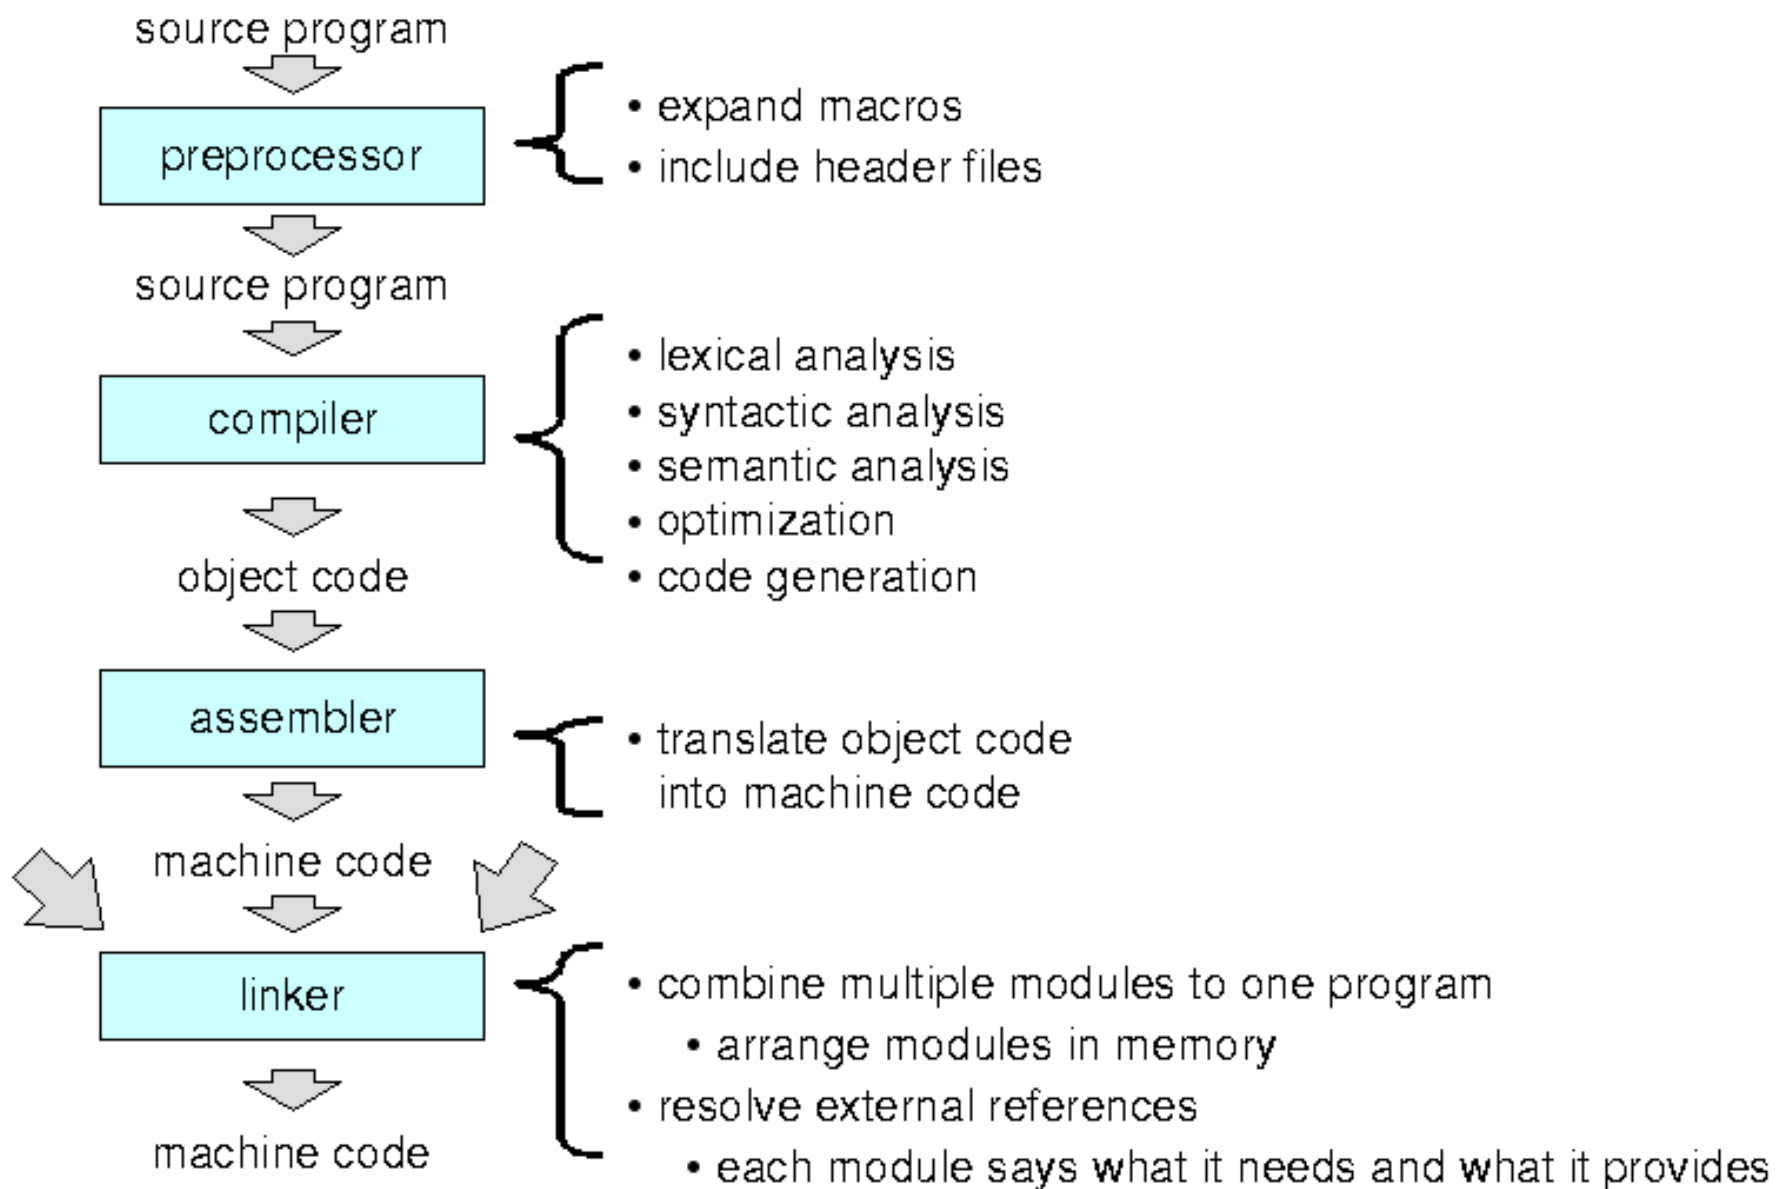
Road map

- Motivation
- Pointers in C

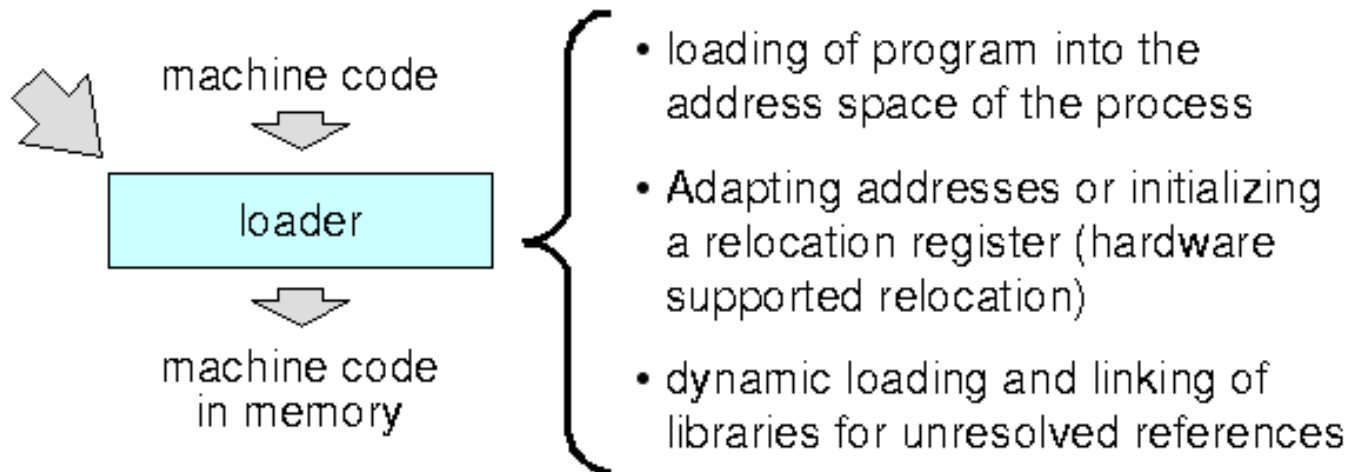
Compilation and memory layout

- Buffer overflows
- Defense
- Summary

Compilation

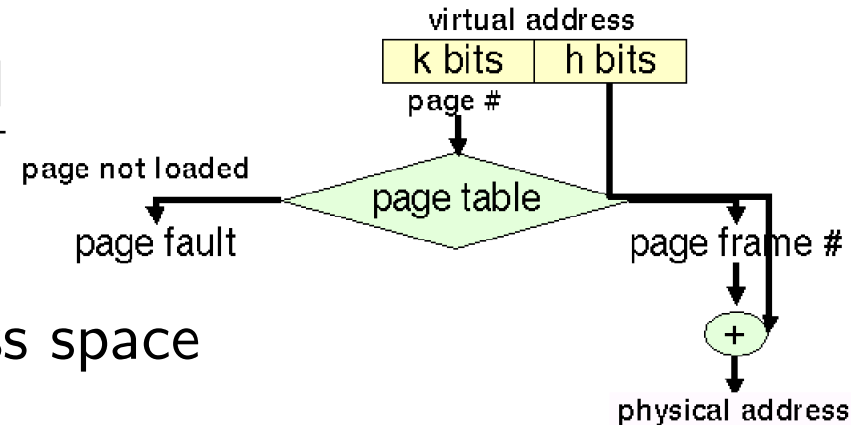


Loading



Recall how (virtual) memory is organized

- Memory named by virtual addresses
- Each process has its own virtual address space
- Protection: process can only access memory in its own space



A program in memory

\$ gdb my-getstring

(gdb) disas main

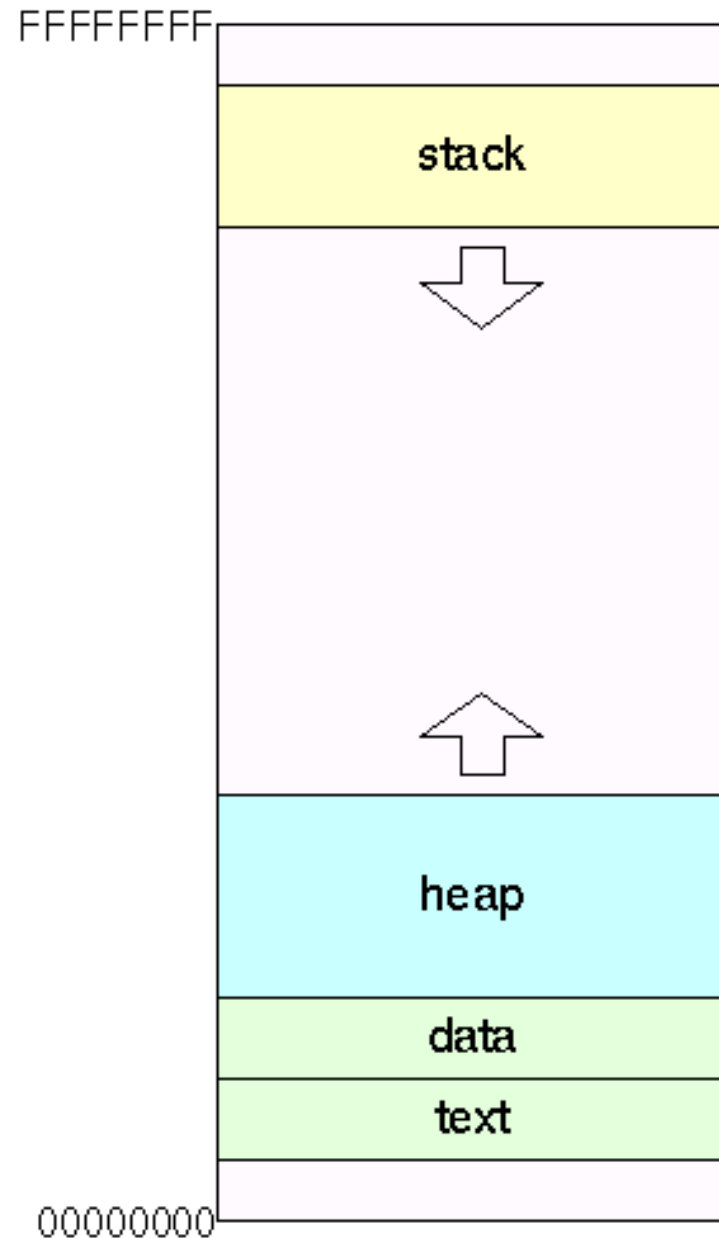
Dump of assembler code for function main:

0x8048440	push	%ebp	0x8048473	incl	(%eax)
0x8048441	mov	%esp,%ebp	0x8048475	jmp	0x804844c <main+ 12>
0x8048443	sub	\$0x18,%esp	0x8048477	nop	
0x8048446	lea	0xffffffff8(%ebp),%eax	0x8048478	mov	0xffffffff0(%ebp),%eax
0x8048449	mov	%eax,0xffffffff0(%ebp)	0x804847b	movb	\$0x0,(%eax)
0x804844c	call	0x80482e8 <getchar>	0x804847e	sub	\$0x8,%esp
0x8048451	mov	%eax,%eax	0x8048481	lea	0xffffffff8(%ebp),%eax
0x8048453	mov	%al,0xffffffff7(%ebp)	0x8048484	push	%eax
0x8048456	mov	0xffffffff7(%ebp),%al	0x8048485	push	\$0x8048508
0x8048459	cmp	\$0xa,%al	0x804848a	call	0x8048328 <printf>
0x804845b	je	0x8048478 <main+56>	0x804848f	add	\$0x10,%esp
0x804845d	cmpl	\$0xff,0xffffffff7(%ebp)	0x8048492	mov	\$0x0,%eax
0x8048461	jne	0x8048468 <main+ 40>	0x8048497	leave	
0x8048463	jmp	0x8048478 <main+56>	0x8048498	ret	
0x8048465	lea	0x0(%esi),%esi	0x8048499	lea	0x0(%esi),%esi
0x8048468	mov	0xffffffff0(%ebp),%edx	0x804849c	nop	
0x804846b	mov	0xffffffff7(%ebp),%al	0x804849d	nop	
0x804846e	mov	%al,(%edx)	0x804849e	nop	
0x8048470	lea	0xffffffff0(%ebp),%eax	0x804849f	nop	

Layout of virtual memory

Linux on Intel x86 family

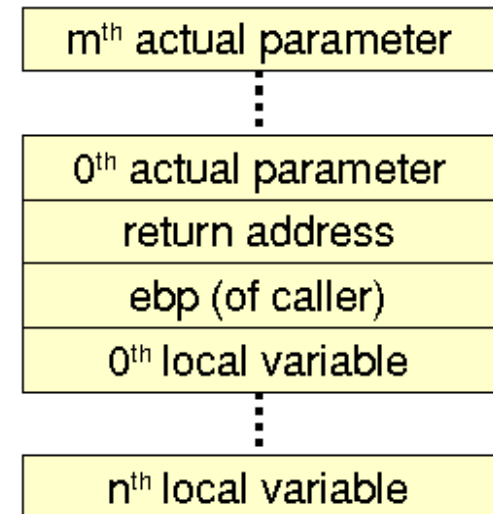
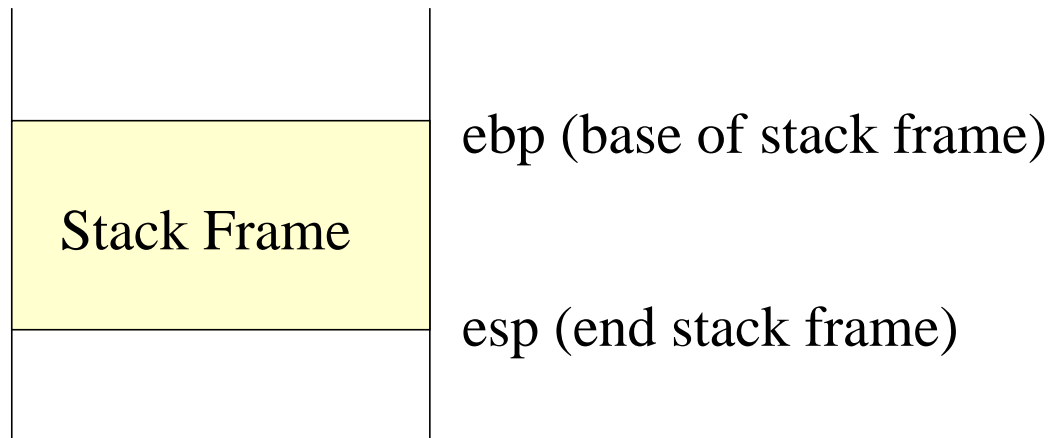
- Stack grows downward
- Stack frame holds
 - ▶ Calling parameters
 - ▶ Local variables for functions
 - ▶ Various addresses
- Heap grows upwards
 - ▶ Dynamically allocated storage generated using `alloc`, `malloc`, ...
- Data: statically allocated storage
- Text: Executable code, read only



Structure of the stack

- Stack grows downwards: one **stack frame** per subroutine called

Higher values



Lower values

- Hardware registers contain:

ebp (extended) base pointer register: start current frame

esp (extended) stack pointer register: top (end) of stack

Another vulnerable program

```
int main (void) {
    char pw[8];
    sprintf (pw, "root-pw");
    if (authenticate(pw) > 0) {
        printf ("[root]$ ...\n");
    } else {
        printf ("Root: incorrect password\n");
    }
    return 0;
}
```

```
#include <stdio.h>
int authenticate (char* pw) {
    int result = 0;
    char buf[8];
    printf ("Root Password: ");

    if (gets(buf) != 0){
        if (strcmp (buf, pw) == 0) {
            result = 1;
        }
    }
    return result;
}
```

gets() non tiene conto della misura di buf
-> fa read da stdin byte by byte

- \$./my-authenticate

Root Password: guess

Root: incorrect password

- \$./my-authenticate

Root Password: root-pw

[root]\$...


Stack dynamics in example

1. start (of main)

base _____ esp

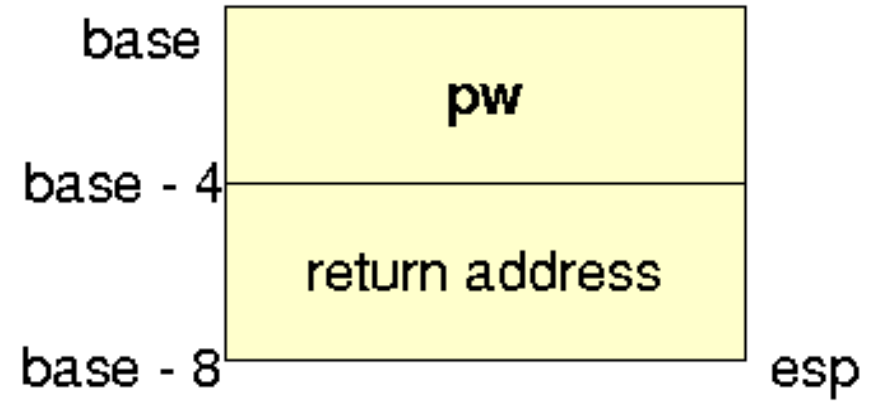
Stack dynamics in example

1. start (of main)
2. push argument pw onto stack



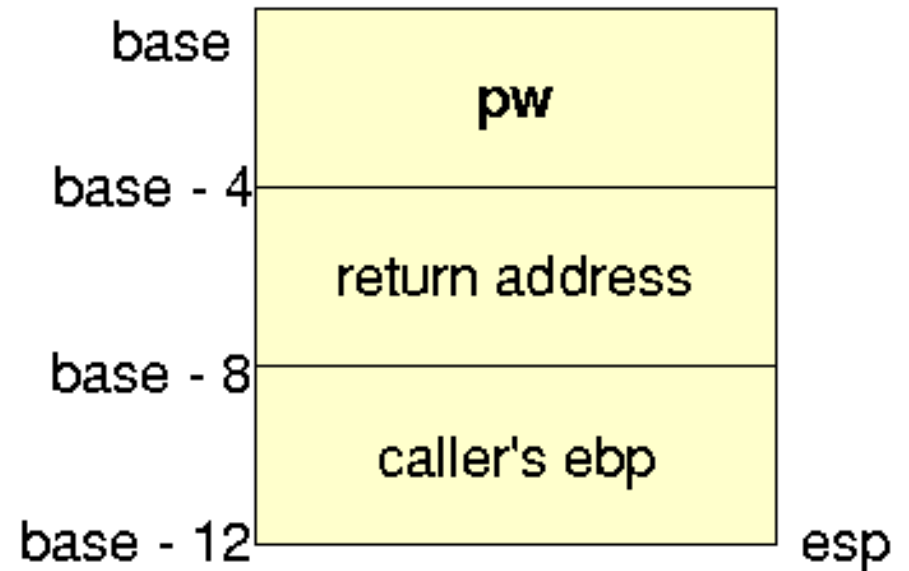
Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)



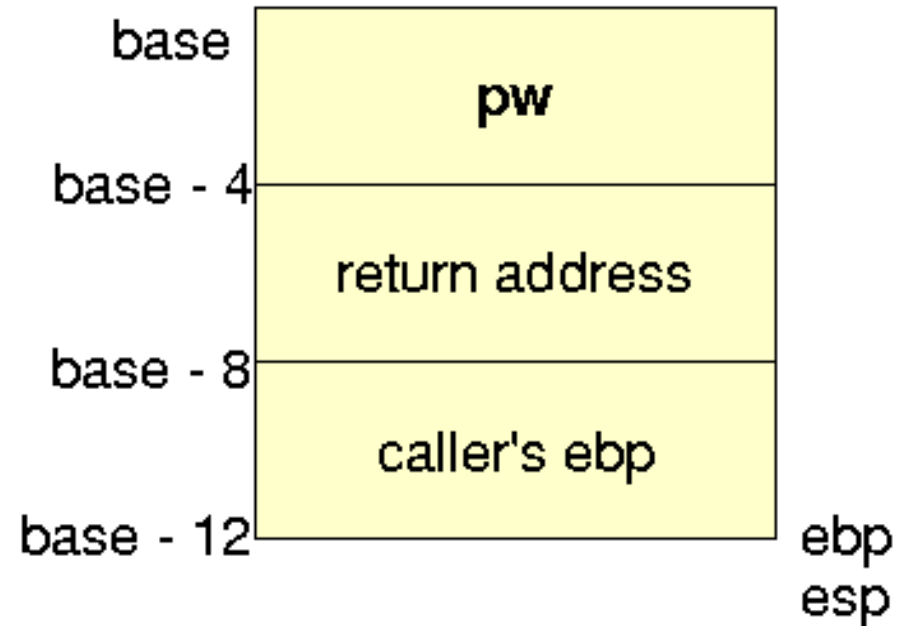
Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)
4. push caller's ebp onto stack



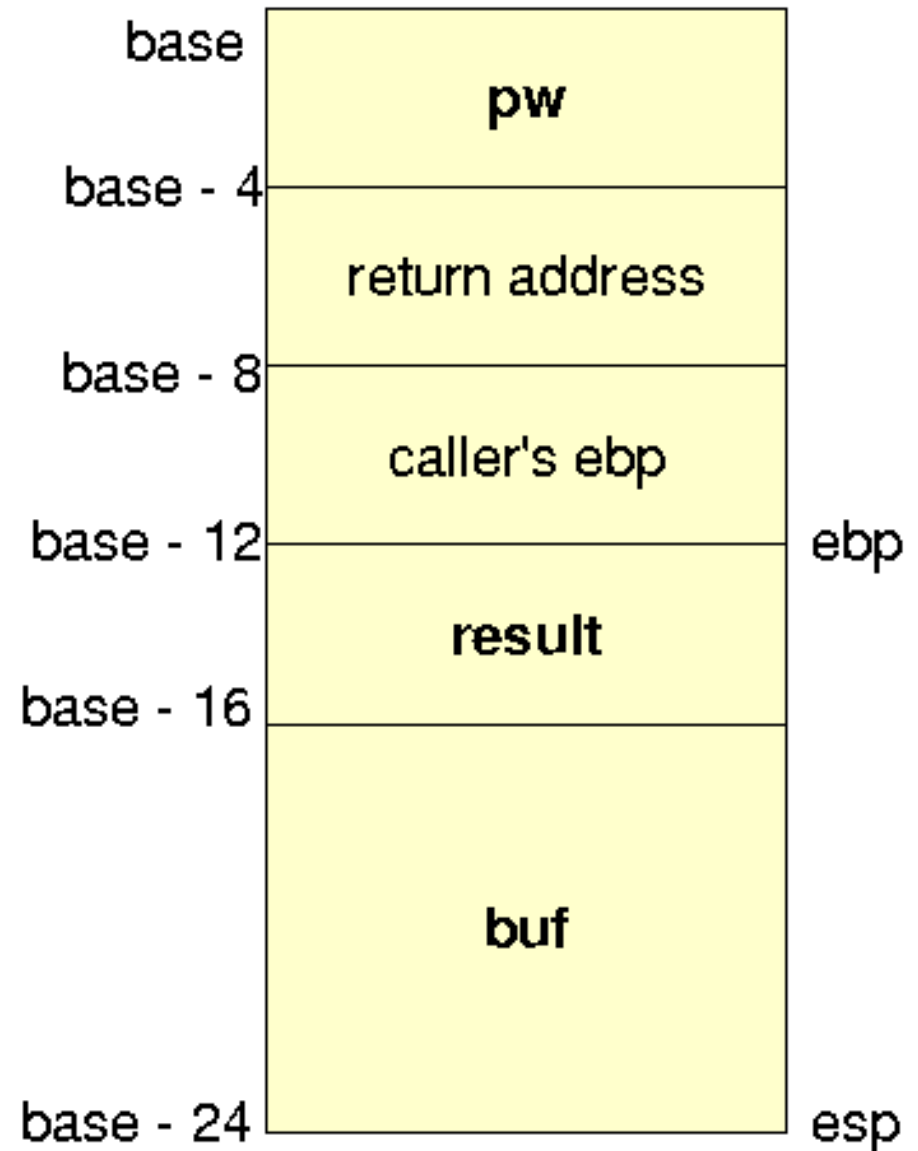
Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)
4. push caller's ebp onto stack
5. copy esp into ebp



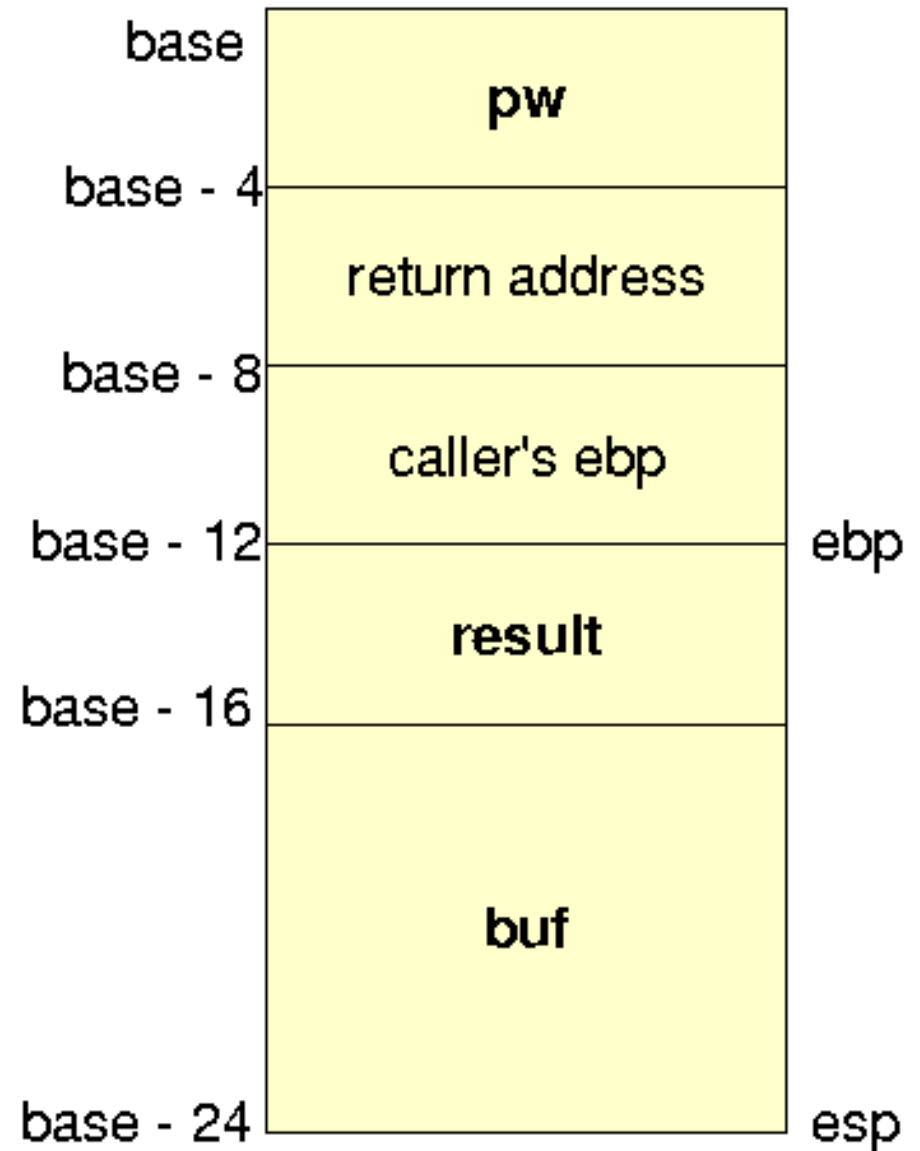
Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)
4. push caller's ebp onto stack
5. copy esp into ebp
6. decrement esp, creating space for local vars



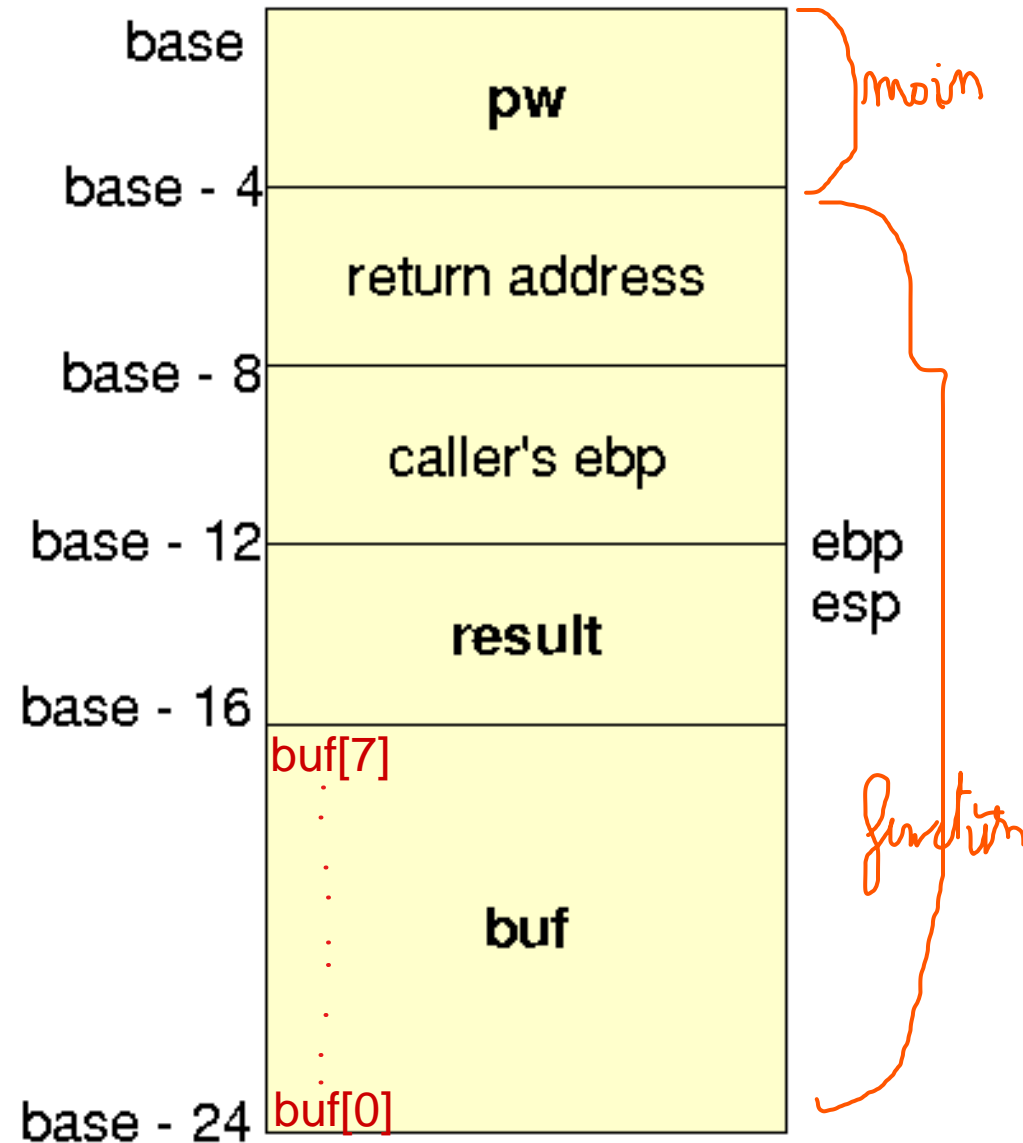
Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)
4. push caller's ebp onto stack
5. copy esp into ebp
6. decrement esp, creating space for local vars
7. Compute authenticate function (stack idle)




Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)
4. push caller's ebp onto stack
5. copy esp into ebp
6. decrement esp, creating space for local vars
7. Compute authenticate function (stack idle)
8. copy caller's ebp into esp



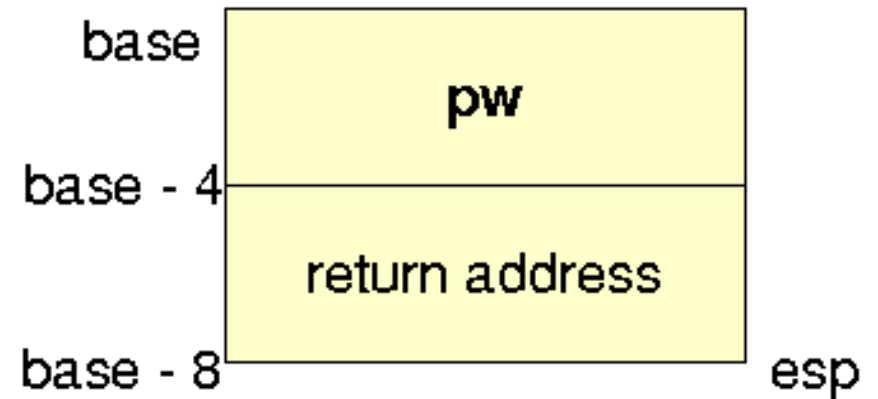
Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)
4. push caller's ebp onto stack
5. copy esp into ebp
6. decrement esp, creating space for local vars
7. Compute authenticate function (stack idle)
8. copy caller's ebp into esp
9. pop vars and caller's ebp from stack



Stack dynamics in example

1. start (of main)
2. push argument pw onto stack
3. call authenticate (pushes return address onto stack)
4. push caller's ebp onto stack
5. copy esp into ebp
6. decrement esp, creating space for local vars
7. Compute authenticate function (stack idle)
8. copy caller's ebp into esp
9. pop vars and caller's ebp from stack
10. return



Road map

- Motivation
- Pointers in C
- Compilation and memory layout

Buffer overflows

- Defense
- Summary

Where is vulnerability in authenticate?

- Some non-problems
 - ▶ Variables `result` and `buf` are initialized prior to use
 - ▶ Programming logic ok.
`result` is 1 iff input equals password supplied by `pw`
- The problem
 - ▶ Implementation of `gets` does not prevent overflow of `buf`
 - ▶ Stack layout

```
#include <stdio.h>
int authenticate (char* pw) {
    int result = 0;
    char buf[8];
    printf ("Root Password: ");

    if (gets(buf) !=0 ){
        if (strcmp (buf, pw) == 0) {
            result = 1;
        }
    }
    return result;
}
```

il problema è che `gets()` legge n bytes finchè non trova un carattere di fine
(quindi anche più di 8 chars)

Example: a long password

- Given a buffer overflow, data from `buf` spills over into `result`
- Can change value of `result`

`result` initialized to 0 and updated only if authentication succeeds

- Works (conceptually) as follows

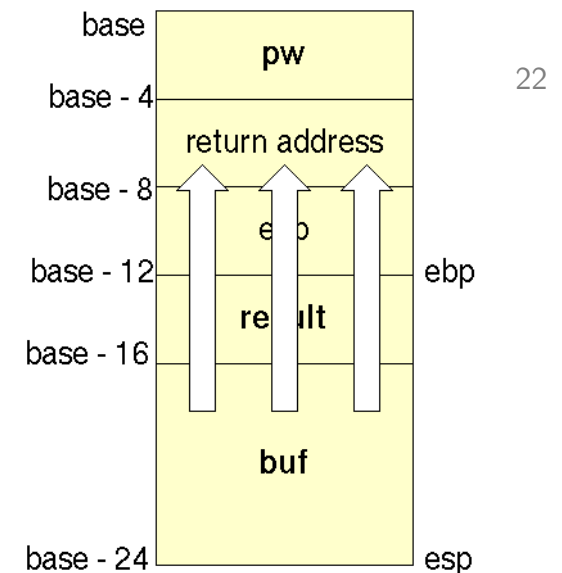
```
$ ./my-authenticate
```

```
Root Password: aaaaaaaa\x01\x00
```

```
[root]$ ...
```

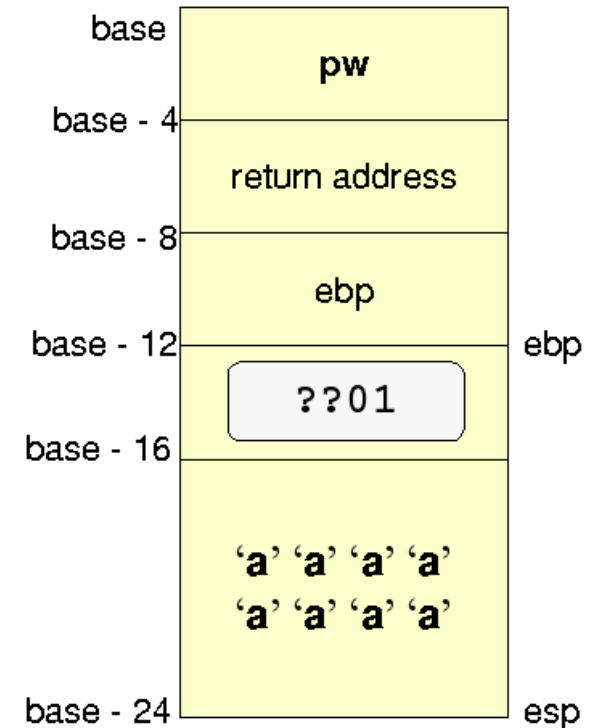
`buf` filled with `aaaaaaaa` and hence `x01` overwrites `result`.

- Violation of data integrity!
 - Variables modified without explicit assignment
 - Thinking purely on the level of C misses this possibility!



Long password (cont.)

- In reality, buffer overflows are a bit trickier
- There could be a gap between `buf` and `result`.
- If input is too long then `ebp` or `return` may be effected



- Entering hexadecimal values `\x01` and `\x00`

Here one can enter input strings using a program via a pipe

Can we program defensively against such errors?

A defense: restoring variables

```
#include <stdio.h>
int authenticate (char* pw) {
    int result = 0;
    char buf[8];
    printf ("Root Password: ");

    if (gets(buf) !=0 ){
        if (strcmp (buf, pw) == 0) {
            result = 1;
        }
    }
    return result;
}
```

```
#include <stdio.h>
int authenticate (char* pw) {
    int result = 0;
    char buf[8];
    printf ("Root Password: ");

    if (gets(buf) !=0 ){
        if (strcmp (buf, pw) == 0) {
            result = 1;
        } else { result = 0; }
    }
    return result;
}
```

- Result is updated after `gets`

Now result is correct even if `buf` overflows into `result`.

- But other dangers still remain!

Would you recommend this defense?

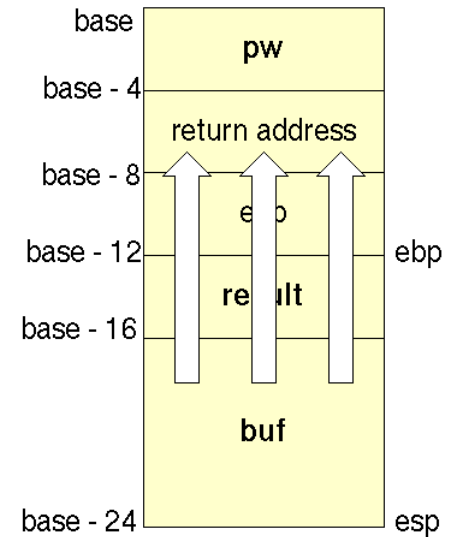
Example: a very long password

- What happens on:

```
$ ./my-authenticate
```

```
Root Password:
```

```
aaaaaaaaaaa\xbb\xbb\xbb\xbb\x11\x47\x11\x47\00+
```



- Overwrites **result**, **ebp**, and **return**


buf: "aaaaaaaa"

result: "aaaa"

stored ebp: bb bb bb bb

return address: 47 11 47 11

*ra milione
pi
bisogna
override
titolo
rete*



- Call of **ret** (when leaving subroutine) pops 47 11 47 11 off stack, starting execution at this memory location

Overflow has destroyed integrity of code-flow!

Exploit code (or where evil doers jump)

- Where would a malicious attacker jump to?

One common target is to code that creates a (root-)shell

- Where in memory does this code go?

- Common approach: place exploit code on the stack

Usually, place within the very buffer that is overflowed

- Return address must then point exactly to the exploit's entry point
 - ▶ Non-trivial in practice.
 - ▶ Trick used of starting exploit code with a “landing zone” of values representing No-op instructions

Exploit code (cont.)

- In above approach, exploit code and landing zone fit into buffer
- Alternatively, attacker places exploit code:
 - ▶ On the stack: into parameters or other local variables
 - ▶ On the heap: into some dynamically allocated memory region
 - ▶ Into environment variables (on stack)
- Another alternative is to abuse existing code
E.g., jump to fragments of the program code or library functions
- Clever abuses have become a kind of hacker sport!
For more on this, see Pincus/Baker article from 2004.

Road map

- Motivation
- Pointers in C
- Compilation and memory layout
- Buffer overflows

Defense


- Summary

Insert a canary

- A **canary** is a **value on the stack whose value is tested before returning.**

```
void function (...) {
  int canary;
  char buf[MAX_SIZE]; // other declarations follow
```

```
  canary = CANARY_VALUE;
  gets(buf);    // or some other dangerous operation
  ...
  if(canary != CANARY VALUE)
    exit
  return; }
```



- A canary is a random value (hard for attacker to guess) or a value composed of different string terminators (CR, LF, Null, -1).
 - Implementations for GCC and Microsoft Visual C++ compilers
- Known attacks exist (see Pincus/Baker)

Defensive programming

- Avoid unsafe library functions, e.g., strcpy, gets, ...
 - ▶ Replace with safe variants, e.g., strncpy, fgets, ...
 - ▶ E.g., replace strcpy(dst,src) with strncpy(dst,src,dst_size-1).
- Always check bound of arrays when iterating over them
- Mechanisms for doing this
 - ▶ Careful programming
 - ▶ Audit teams
 - ▶ Use grep or more sophisticated tools
- Limitations: error prone and audits are time consuming

Defensive programming (cont.)

BUGS

Never use `gets()`. Because it is impossible to tell without knowing the data in advance how many characters `gets()` will read, and because `gets()` will continue to store characters past the end of the buffer, it is extremely dangerous to use. It has been used to break computer security. Use `fgets()` instead.

— From `gets(3)` manual page

Automatic array bounds checking

- **Approach:** compiler automatically adds an explicit check to each access to an array

- Checks inserted during code generation

Example: GCC enhancements of Jones&Kelly.

- Drawbacks
 - ▶ It can be difficult to determine the bounds of an array
 - ▶ Loss of performance can be substantial, e.g., factor 10-30!
 - ▶ Some compilers only check explicit array references, like `buf[n]`, but not pointer references, like `*(buf+n)`

Avoid C (++)

- Use a language that is type safe
 - ▶ Length of an array is part of its type
 - ▶ Assigning contents of a buffer to a smaller buffer is a type error
 - ▶ **Examples:** Pascal, Java, or ML
- Problems and limitations
 - ▶ Often the choice of programming language is not yours
 - ▶ Bugs still possible if run-time environment is programmed in an unsafe language
 - Example:** several buffer overflows vulnerabilities in implementations of the JAVA virtual machine
 - ▶ C(++) has its advantages, in particular for writing efficient programs “close to the machine”.

Avoid buffers on stack

- Avoid declaration of local array variables in functions.
- Instead, use heap storage, e.g., allocate space with `malloc()`.
- As return address is on the stack, it can not be overwritten by a buffer-overflow on the heap.
- Does this solve all our worries?

Avoid buffers on stack

- Avoid declaration of local array variables in functions.
- Instead, use heap storage, e.g., allocate space with `malloc()`.
- As return address is on the stack, it can not be overwritten by a buffer-overflow on the heap.
- Does this solve all our worries? No!
 - ▶ Heap overflows are also a real problem (not covered here)
 - ▶ While they have no effect on control-flow integrity, they can violate data integrity.

Non-Executable Buffers

- Mark stack [or heap] as being non-executable.
 - ⇒ attacker cannot run exploit stored in buffers on stack [heap].
- Mechanisms
 - ▶ Extend OS with a register storing maximal executable address
 - ▶ Alternatively, tag pages as (non)executable in the page table
- Problems and limitations
 - ▶ Attacker can still execute code in the text segment
 - ▶ Attacker can still violate data integrity
 - ▶ Sometimes too restrictive

Unix signal handlers usually execute on the current process stack. This lets the signal handler return to the point that execution was interrupted in the process.

Road map

- Motivation
- Pointers in C
- Compilation and memory layout
- Buffer overflows
- Defense

 **Summary**

Conclusions and lessons learned

- Arrays in C can be over-written.
- The result is that data and control flow can be altered in ways not described by program itself.
- This is a massive problem and has been so for many years!
- Defense is paramount!
 - ▶ Program defensively: only use or write functions that do bounds checking.
 - ▶ Carefully weigh pros and cons of programming language used. Do advantages of C outweigh its disadvantages?
 - ▶ Compiler/hardware support for preventing overflows is available. It helps, but be aware of the limitations and overhead.

Literature

- John Viega, Gary McGraw. Building Secure Software. Addison-Wesley, 2002.
- Crispin Cowan, Perry Wagle, Calton Pu, Steve Beattie, Jonathan Walpole. *Buffer Overflows: Attacks and Defences for the Vulnerability of the Decade*. DARPA Information Survivability Conference and Exposition, 2000.
- Jonathan Pincus, Brandon Baker. *Beyond Stack Smashing: Recent Advances in Exploiting Buffer Overruns*, IEEE Security & Privacy, 2004.
- AlephOne. Smashing the Stack for Fun and Profit. 1996